



## A Calculus for Communicating Systems?

## A Calculus for Communicating Systems (CCS)

$  \begin{array}{l}  P, Q ::= \emptyset \\    \text{ in } a(b).P \\    \text{ out } a\langle b \rangle.P \\    !P \\    \text{ new } a;P \\    (P \mid Q) \\    P + Q \\    \text{ if } (\text{cond}) \\  \quad \text{then } P  \end{array}  $	$  \begin{array}{l}  P, Q ::= \emptyset \\    a(b).P \\    a\langle b \rangle.P \\    !P \\    \nu a;P \\    (P \mid Q) \\    P + Q \\    [x = y]P  \end{array}  $
--	--

## Semantics

$$\begin{array}{c}
 \text{out } a\langle b \rangle \xrightarrow{\bar{a}\langle b \rangle} 0 \qquad \text{in } a(x);P \xrightarrow{a(x)} P \\
 \frac{P \xrightarrow{\bar{a}\langle b \rangle} P' \quad Q \xrightarrow{a(x)} Q'}{P \mid Q \xrightarrow{\bar{a}\langle b \rangle} P' \mid Q'[b/x]} \\
 \\
 \frac{P \xrightarrow{\alpha} P'}{P \mid Q \xrightarrow{\alpha} P' \mid Q} \qquad \frac{Q \xrightarrow{\alpha} Q'}{P \mid Q \xrightarrow{\alpha} P \mid Q'} \\
 \\
 !P \rightarrow P \mid !P \qquad \frac{P \xrightarrow{\alpha} P' \quad a \notin \alpha}{\text{new } a.P \xrightarrow{\alpha} \text{new } a.P'} \\
 \\
 \frac{P \equiv Q \quad Q \rightarrow Q'}{P \rightarrow P} \qquad \frac{P \rightarrow P' \quad P' \rightarrow P'' \quad P'' \rightarrow P'''}{P \rightarrow P'''}
 \end{array}$$

## Pi-calculus

- The pi-calculus extends CCS by letting channel names be rewritten.

<p>Client:</p> $  \begin{array}{l}  \nu c . \overline{\text{server}}(\text{finger}, c) \\    c(x).\text{print}(x)  \end{array}  $ <p>Server:</p> $  \begin{array}{l}  \text{server}(\text{service}, \text{reply}).\overline{\text{service}}(\text{reply}) \\    \text{finger}(\text{reply}).\overline{\text{reply}}(\text{users}) \\    \text{time}(\text{reply}).\overline{\text{reply}}(\text{now})  \end{array}  $	<p>← ask server</p> <p>← print response</p> <p>← inet daemon</p> <p>← finger daemon</p> <p>← time daemon</p>
---	--

- Most models don't use this, and are really CCS!

## Weak bi-simulation

- Two systems are weakly bi-similar if:
  - every input and output of one systems is matched with the same input or output and any number of internal transitions in the other system
  - and the resulting processes are also bi-similar.
- Weak bi-simulation defines the observation power of the attacker.
- Using bi-simulation, rather than trace equivalence gives the attacker the power to detected states where a possible action has been ruled out.

## Weak bi-simulation


- A definition of looks the same:

**Definition 1 (Weakly bi-similar).** Processes  $P$  and  $Q$  are weakly bi-similar if there exists an equivalence relation  $\approx$  such that  $P \approx Q$  and for all  $P_1$  and  $Q_1$  such that  $P_1 \approx Q_1$ , if  $P_1 \xrightarrow{\alpha} P'_1$  then:

- if  $\alpha$  is an output or an internal action there exists a process  $Q'_1$  such that  $Q_1 \xrightarrow{\alpha} Q'_1$  and  $P'_1 \approx Q'_1$ .
- if  $\alpha$  is an input action, i.e.,  $\alpha = a(x)$ , then for all names  $b$ , there exists a process  $Q'_1$  such that  $Q_1 \xrightarrow{\alpha} Q'_1$  and  $P_1[b/x] \approx Q'_1[b/x]$ .



## Peer-to-Peer File-Sharing




In newer networks peers record the IP address of other peers.

A searcher sends a request to all of its "neighbours".

This is forwarded to all of their neighbours, up to a fixed hops.

## Peer-to-Peer File-Sharing




The search request includes A's IP address.

Any peer with the requested file contacts A directly.

Peer "A" may then request the file.

## Peer-to-Peer File-Sharing




No anonymity from peers inside the network:

The search message gives the searcher's IP address and name of the files they are looking for.

By requesting a file, you can find out the IP address of all peers that are offering the file.

## MUTE: Search




The search takes place as before, but this time the message uses its pseudo ID as the "from ID".

Each peer builds a routing table by records the ID and the connection.

A probabilistic time-to-live counter limits the search.

## MUTE: Reply



If B wants to reply it sends a message to A's pseudo ID.

This message is routed using the ad-hoc routing table.

The route to B is also recorded

## Anonymity Provided by MUTE

It should be impossible for an attacker who

- is acting as any number of peers,
- and can observe any number of connections,
- but does not know what the network looks like,

i.e., cannot be sure that they have surrounded a peer to be able to link any IP address and pseudo ID with an arbitrarily high probability.


## The Pi-calculus Model of MUTE

- We make pi-calculus models of
  - an "innocent" peer:  $I$ , that only forward messages.
  - A "guilty" peer:  $G$ , that forwards messages and replies to requests for a file.
  - And an environment:  $E$ ,
- These processes are parameterised on the communication channels that run between them...
- and the choices the peer makes at start up.

## The Pi-calculus Model of MUTE

- We can use these processes as building blocks to make any network.

e.g.  $I(c_1, c_2, c_3) \mid G(c_2, c_4) \mid G(c_3, c_5) \mid E(c_4, c_5)$   
 e.g.  $\text{new } c_2, c_3, c_4, c_5; I(c_1, c_2, c_3) \mid G(c_2, c_4) \mid G(c_3, c_5) \mid E(c_4, c_5)$



## Model of an Innocent Peer.

```
InnocentPeer( (c1,co1), (c2,co2), for, ttl )
= new my_id,
  !(in c1; (kind,to_id,from_id,phase,counter);
    out <memory, (from_id,co1)>;
    | if [kind = search] then Forward
    | if [kind = reply] then Reply
    ... )
  | ! (out<co1, (search,_,my_id,1,ttl) ) }


Reply = if (to_id = my_id) then (stop)
| if (to_id ≠ my_id) then
  new loop; out loop;
  !( in loop;
    in memory(x, channel);
    out memory(x, channel);
    if (x ≠ to_id) then out loop;
    | if (x = to_id) then
      out channel(kind, to_id, from_id, phase, counter)
```

## Correctness

We want to show that the environment provides cover for a guilty peer:

i.e., for all guilty peers  $G$  and environments  $\text{Env}$  there exists an innocent peer  $I$  and environment  $\text{Env}'$  such that

$$G \mid E \approx I \mid E'$$



## Checking the Model

We check that:

$\text{new } c_2.( \text{Guilty\_Peer}(c_1, c_2, p_g) \mid \text{Env}(c_2, e_1) )$   
 $\approx \text{new } c_2.( \text{Innocent\_Peer}(c_1, c_2, p_i) \mid \text{Env}(c_2, e_2) )$

This is small enough to do by hand  
 ~ 24 states to check.


## Failure of Bisimulation

A state that cannot be matched:

- receive a search message from the attacker,
- sent a reply from the guilty peer / the environment,
- receive a well-formed search message from the guilty peer's ID,
- receive a reply message "from" the guilty peers ID.
- Then the innocent network can perform .
  - a receive the same reply back
- The guilty network cannot perform this action.


## What Happened?

Guilty:



- The guilty network shares its file and makes its ID public.
- The innocent network matches this with a file from the environment.


Innocent:



- The attacker makes a search message using this ID as the **from** address and sends it back.


## What Happened?

Guilty:




- A record of that ID and the connection is added to peers routing table.
- This is an "incorrect" entry in the routing table.
- It routes messages away from their real owner and to the attacker.

Innocent:




## What Happened?

Guilty:



- Now the attacker sends a reply message **to** that ID.
- The innocent peer may send it back.

Innocent:



- Whereas, the guilty peer will always accept the message.

## Description of the Real Attack

- The attacker tries to steal an address by sending many fake message.
- The attacker then sends 50+ reply messages, use the target ID.
- If the reply is not returned to the attacker then there is a high probability that it belongs to the neighbour.

## Fixing MUTE

- The final solution was for peers to generate RSA signature keys “kS” and an authentication keys “kA”.
- Message header has the form:  

$$( \#(kA), \text{from ID}, \text{message ID-time\_stamp}, \text{FLAGS:(S}_{kS}(\text{messageID-time\_stamp}), kA) )$$
- This fix was added to the latest version of MUTE.

## Conclusion

- Modelling protocols in process calculi is a great way to understand them
- We can test anonymity using bisimulation:

For all  $p_g, e_1$  there exists  $p_i, e_2$  s. t.  
 $\text{Guilty\_Peer}(p_g) \mid \text{Env}(e_1) \approx \text{Innocent\_Peer}(p_i) \mid \text{Env}(e_2)$

- 500,000 people downloaded and used the fixed version (before they all gave up and used BitTorrent instead).

## Summary

- A formal language lets us model and analyse complex protocols.
- This is an excellent tool for protocol analysis.
- If you ever analyse a protocol, or a similar complex system give this a try.
- BUT, this is all ancient history. Today we have excellent fully automatic tools to do the analysis for us.
- These tools will be the subject of tomorrow's lecture.